

```

## LINK FOR DATASET: https://drive.google.com/drive/folders/1ks-KSDBHB1FYsXtt_T_HclX-7dxBL7Nw?usp=sharing
import pandas as pd
import glob
import os
import numpy as np
import warnings
from google.colab import drive

# Suppress warnings for cleaner output
warnings.filterwarnings('ignore')

# --- I. Configuration and Column Mapping ---

# The permanent and correct mapping based on your raw data sample
CANONICAL_COL_MAP = {
    'DAY_OF_WEEK': 'DayOfWeek',
    'CRS_DEP_TIME': 'CRSDepTime',
    'OP_UNIQUE_CARRIER': 'Reporting_Airline',
    'DEST_AIRPORT_ID': 'DestAirportID',
    'ORIGIN_AIRPORT_ID': 'OriginAirportID',
    'DISTANCE': 'Distance',
    'DEP_DELAY': 'DEP_DELAY',
    'CANCELLED': 'CANCELLED',
    'FL_DATE': 'FL_DATE',
    'LATE_AIRCRAFT_DELAY': 'LateAircraftDelay' # Crucial for propagation feature
}

RAW_REQUIRED_COLS = list(CANONICAL_COL_MAP.keys())

# DOT IDs for the Five Strategic Hubs: ORD, MDW, MKE, DTW, and MSP
AIRPORT_IDS = [11298, 10821, 13244, 11433, 13487]
data_path = '/content/drive/MyDrive/CS441/Final Project/Monthly Raw Data'
ORIGIN_COL = 'ORIGIN_AIRPORT_ID'

# -----
print("--- Starting Data Assembly and Filtering ---")
drive.mount('/content/drive')

# 1. Concatenate Files (Recursive Search)
# Note: Added '**' and 'recursive=True' to find files in subfolders, which was a fix earlier.
file_pattern = os.path.join(data_path, '**', '*.csv')
all_files = glob.glob(file_pattern, recursive=True)

if not all_files:
    print(f"FATAL ERROR: No CSV files found in {data_path}. Check the path.")
    exit()

print(f"Found {len(all_files)} files. Starting concatenation...")

try:
    # Use usecols to only load the columns we need, saving memory
    list_of_dfs = [pd.read_csv(f, usecols=RAW_REQUIRED_COLS, low_memory=False) for f in all_files]
    df = pd.concat(list_of_dfs, ignore_index=True)
except Exception as e:
    print(f>An error occurred during reading or concatenation: {e}")
    exit()

initial_total_rows = df.shape[0]
print(f"\nInitial concatenated dataset size: {initial_total_rows} rows.")

# Apply Column Mapping and Deduplication
df.rename(columns=CANONICAL_COL_MAP, inplace=True)
DEDUP_COLS = ['FL_DATE', 'Reporting_Airline', 'OriginAirportID', 'CRSDepTime']
df.drop_duplicates(subset=DEDUP_COLS, inplace=True, keep='first')
rows_removed_by_dedup = initial_total_rows - df.shape[0]
print(f"Removed {rows_removed_by_dedup} duplicate rows.")

# 2. Filter to Five Hubs and Clean Core Data
df['OriginAirportID'] = pd.to_numeric(df['OriginAirportID'], errors='coerce').fillna(0).astype('Int64')
df_final = df[df['OriginAirportID'].isin(AIRPORT_IDS)].copy()

core_cols_for_check = ['DEP_DELAY', 'CANCELLED', 'CRSDepTime', 'Reporting_Airline']
initial_rows = df_final.shape[0]
df_final.dropna(subset=core_cols_for_check, inplace=True)

```

```

df = df_final # Use 'df' for the final working DataFrame

final_rows = df.shape[0]
print(f"Filtered and cleaned dataset size: {final_rows:,} rows.")
print("--- Data Assembly Complete. Starting Feature Engineering ---")

# --- II. Feature Engineering ---

# Feature 1: The Target Variable (Y)
print("\n--- Feature 1: Target Variable ---")
df['TARGET_CLASS'] = 0
df.loc[(df['DEP_DELAY'] > 15) & (df['CANCELLED'] == 0), 'TARGET_CLASS'] = 1 # Significant Delay
df.loc[df['CANCELLED'] == 1, 'TARGET_CLASS'] = 2 # Cancellation
print(f"Target Class Distribution:\n{df['TARGET_CLASS'].value_counts(normalize=True).mul(100).round(2).astype(str)} +")
df.drop(columns=['DEP_DELAY', 'CANCELLED'], inplace=True)

# Feature 2: Cyclical Encoding for Time
print("\n--- Feature 2: Cyclical Time Encoding ---")
df['Time_of_Day_Minutes'] = df['CRSDepTime'] // 100 * 60 + df['CRSDepTime'] % 100
MAX_MINUTES = 24 * 60

# Sin/Cos transformation
df['DepTime_sin'] = np.sin(2 * np.pi * df['Time_of_Day_Minutes'] / MAX_MINUTES)
df['DepTime_cos'] = np.cos(2 * np.pi * df['Time_of_Day_Minutes'] / MAX_MINUTES)
df.drop(columns=['CRSDepTime', 'Time_of_Day_Minutes'], inplace=True)
print("Created DepTime_sin and DepTime_cos features.")

# Feature 3: The Lagged Delay Propagation Feature
print("\n--- Feature 3: Lagged Delay Propagation (Extraordinary Feature) ---")

def calculate_lagged_mean(group, column, window_size=50):
    """Calculates the rolling mean for a column, shifted by 1."""
    return group[column].shift(1).rolling(window=window_size, min_periods=1).mean()

# Prepare data: Ensure correct data types and chronological sort for the rolling calculation
df['FL_DATE'] = pd.to_datetime(df['FL_DATE'])
# Sort by Hub, then Airline, then Chronologically (Date and Time proxy)
df.sort_values(by=['OriginAirportID', 'Reporting_Airline', 'FL_DATE', 'DepTime_sin'], inplace=True)

# Lagged_Late_Aircraft: Average LateAircraftDelay for the *previous 50 flights* by this airline at this hub.
df['Lagged_Late_Aircraft'] = df.groupby(['OriginAirportID', 'Reporting_Airline']) \
    .apply(calculate_lagged_mean, 'LateAircraftDelay', 50) \
    .reset_index(level=[0,1], drop=True)

# Lagged_Delay_Mean: Average TARGET_CLASS for the *previous 50 flights* by this airline at this hub.
df['Lagged_Delay_Mean'] = df.groupby(['OriginAirportID', 'Reporting_Airline']) \
    .apply(calculate_lagged_mean, 'TARGET_CLASS', 50) \
    .reset_index(level=[0,1], drop=True)

df['Lagged_Late_Aircraft'].fillna(0, inplace=True)
df['Lagged_Delay_Mean'].fillna(0, inplace=True)
df.drop(columns=['LateAircraftDelay'], inplace=True)
print("Created Lagged_Late_Aircraft and Lagged_Delay_Mean.")

# 4. Final Save (Feature Engineered Data)
MASTER_FE_FILE_PATH = '/content/drive/MyDrive/CS441/Final Project/Five_Hub_FE_Master_Data.csv'
df.to_csv(MASTER_FE_FILE_PATH, index=False)
print(f"\nFINAL FEATURE-ENGINEERED dataset saved to: {MASTER_FE_FILE_PATH}")
print("\n--- NEXT STEP: Categorical Encoding and XGBoost Model Training ---")

```

--- Starting Data Assembly and Filtering ---  
 Mounted at /content/drive  
 Found 46 files. Starting concatenation...

Initial concatenated dataset size: 2,771,870 rows.  
 Removed 591,317 duplicate rows.  
 Filtered and cleaned dataset size: 111,774 rows.  
 --- Data Assembly Complete. Starting Feature Engineering ---

--- Feature 1: Target Variable ---  
 Target Class Distribution:  
 TARGET\_CLASS  
 0 80.26%  
 1 19.7%  
 2 0.04%

```
Name: proportion, dtype: object  
--- Feature 2: Cyclical Time Encoding ---  
Created DepTime_sin and DepTime_cos features.  
--- Feature 3: Lagged Delay Propagation (Extraordinary Feature) ---  
Created Lagged_Late_Aircraft and Lagged_Delay_Mean.  
INFO: FINAL FEATURE-ENGINEERED dataset saved to: /content/drive/MyDrive/CS441/Final Project/Five_Hub_FE_Master_Data.csv  
--- NEXT STEP: Categorical Encoding and XGBoost Model Training ---
```

```
# Install the category_encoders library  
!pip install category_encoders
```

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from category_encoders import TargetEncoder
import os
import time

# Check if the file exists before attempting to load
if not os.path.exists(MASTER_FE_FILE_PATH):
    print(f"FATAL ERROR: Feature Engineered file not found at {MASTER_FE_FILE_PATH}.")
    print("Please confirm the file path in your Google Drive and try again.")
    # Exit or stop execution here if the file is missing
    # return
else:
    print(f"Loading feature-engineered data from: {MASTER_FE_FILE_PATH}")
    df = pd.read_csv(MASTER_FE_FILE_PATH)

    print(" --- Starting Categorical Encoding and Data Split ---")

    # --- 1. Define Features and Target ---
    TARGET = 'TARGET_CLASS'
    # Drop target and the date column (FL_DATE) as it's not a direct feature
    FEATURES = df.drop(columns=[TARGET, 'FL_DATE']).columns.tolist()

    # Identify Categorical Columns for Encoding
    CAT_COLS = ['OriginAirportID', 'DestAirportID', 'Reporting_Airline']

    # Ensure categorical columns are treated as strings for the encoder
    for col in CAT_COLS:
        df[col] = df[col].astype(str)

    # --- 2. Data Split (Crucial for Target Encoding) ---
    X = df[FEATURES]
    y = df[TARGET]

    # Split data into training and testing sets (80/20 split)
    # Stratify ensures the rare classes (1 and 2) are distributed evenly in both sets.
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.20, random_state=42, stratify=y
    )

    print(f"Data split: Training set size: {X_train.shape[0]}, Testing set size: {X_test.shape[0]}")
```

```
# --- 3. Target Encoding (Applied only to training data) ---
# Target Encoding is necessary for high-cardinality features.
encoder = TargetEncoder(cols=CAT_COLS)

# Fit the encoder ONLY on the training data (y_train) to prevent data leakage.
encoder.fit(X_train, y_train)

# Transform both the training and testing sets.
X_train_encoded = encoder.transform(X_train)
X_test_encoded = encoder.transform(X_test)

print("Applied Target Encoding to Origin, Destination, and Airline features.")

# --- 4. Final Data Preparation ---
y_train_int = y_train.astype(int)
y_test_int = y_test.astype(int)

# Save the encoded dataframes for the next step (XGBoost)
# The variables X_train_encoded, X_test_encoded, y_train_int, y_test_int are now ready.
print("--- Categorical Encoding and Data Split Complete. Ready for Modeling! ---")
print(f"Features ready for XGBoost: {X_train_encoded.columns.tolist()}")
```

Loading feature-engineered data from: /content/drive/MyDrive/CS441/Final Project/Five\_Hub\_FE\_Master\_Data.csv  
--- Starting Categorical Encoding and Data Split ---  
Data split: Training set size: 89,419, Testing set size: 22,355  
Applied Target Encoding to Origin, Destination, and Airline features.  
--- Categorical Encoding and Data Split Complete. Ready for Modeling! ---  
Features ready for XGBoost: ['DayOfWeek', 'Reporting\_Airline', 'OriginAirportID', 'DestAirportID', 'Distance', 'DepTi

```
import xgboost as xgb
from xgboost import XGBClassifier
from sklearn.metrics import f1_score, confusion_matrix, classification_report
import pandas as pd
import numpy as np
import warnings

warnings.filterwarnings('ignore')
print("--- Starting 5-Model Hyperparameter Search (Targeting Extraordinary Score) ---")

# --- DATA SIZE SUMMARY (FOR REPORT) ---
# NOTE: The variables X_train_encoded and X_test_encoded are available from the previous step.
# We explicitly report the size here as requested.
print("\n Data Split Summary for Report:")
print(f"    Training Set Size (X_train): {X_train_encoded.shape[0]:,} rows")
print(f"    Testing Set Size (X_test):   {X_test_encoded.shape[0]:,} rows")
print(f"    Total Samples Used:        {X_train_encoded.shape[0] + X_test_encoded.shape[0]:,} rows")
print("-----")

# Define five distinct parameter configurations for a comprehensive search
param_sets = {
    "Low_Complexity": {
        "n_estimators": 100,
        "max_depth": 4,
        "learning_rate": 0.15
    },
    "Optimal_Balance": {
        "n_estimators": 150,
        "max_depth": 6,
        "learning_rate": 0.1
    },
    "High_Complexity": {
        "n_estimators": 250,
        "max_depth": 8,
        "learning_rate": 0.05
    },
    "Low_Learning_Rate": {
        "n_estimators": 200,
        "max_depth": 7,
        "learning_rate": 0.02
    },
    "Aggressive_Learning": {
        "n_estimators": 120,
        "max_depth": 5,
        "learning_rate": 0.25
    }
}
```

```

        }

    }

    # Calculate Class Weights
    class_counts = y_train_int.value_counts().sort_index()
    total_samples = len(y_train_int)
    scale_factor = total_samples / (len(class_counts) * class_counts)
    sample_weights = y_train_int.map(scale_factor).values

    results = {}

    # Iterate through all parameter sets
    for name, params in param_sets.items():
        print(f"\nTraining Model: {name}...")

        # CRITICAL FIX: eval_metric is passed during initialization for this XGBoost version
        xgb_model = XGBClassifier(
            objective='multi:softmax',
            num_class=3,
            random_state=42,
            tree_method='hist',
            subsample=0.8,
            colsample_bytree=0.8,
            eval_metric='mlogloss',
            **params
        )

        # --- START TIME MEASUREMENT ---
        start_time = time.time()
        xgb_model.fit(
            X_train_encoded,
            y_train_int,
            sample_weight=sample_weights,
            verbose=False
        )
        # --- END TIME MEASUREMENT ---
        end_time = time.time()
        training_duration = end_time - start_time

        y_pred = xgb_model.predict(X_test_encoded)

        # Use Weighted F1-Score as the primary evaluation metric
        f1_weighted = f1_score(y_test_int, y_pred, average='weighted', zero_division=0)

        results[name] = {
            'F1_Score': f1_weighted,
            'Model': xgb_model,
            'Predictions': y_pred
        }

        print(f"  {name} Weighted F1-Score: {f1_weighted:.4f} | Time: {training_duration:.2f}s")

    # --- 2. Select the Best Model ---
    best_model_name = max(results, key=lambda k: results[k]['F1_Score'])
    best_result = results[best_model_name]
    best_f1 = best_result['F1_Score']
    best_predictions = best_result['Predictions']
    y_test = y_test_int

    print("\n--- Model Selection Complete ---")
    print(f"💡 Best Model Selected: {best_model_name} (F1-Score: {best_f1:.4f})")
    print("-----")

    # --- 3. Final Evaluation of the Best Model ---

    # Detailed Classification Report
    print("\n--- Detailed Classification Report (Labels: 0=OnTime, 1=Delay, 2=Cancel) ---")
    print(classification_report(y_test, best_predictions, target_names=['OnTime', 'Delay', 'Cancel'], zero_division=0))

    # Confusion Matrix (Crucial for visualization and analysis)
    conf_mat = confusion_matrix(y_test, best_predictions)
    print("\n--- Confusion Matrix ---")
    print(conf_mat)

    # Prepare the Hyperparameter Comparison Table for the report

```

```

comparison_df = pd.DataFrame({
    'Model': list(param_sets.keys()),
    'n_estimators': [p['n_estimators'] for p in param_sets.values()],
    'max_depth': [p['max_depth'] for p in param_sets.values()],
    'learning_rate': [p['learning_rate'] for p in param_sets.values()],
    'Weighted F1-Score': [results[name]['F1_Score'] for name in param_sets.keys()]
}).round(4)

print("\n--- Hyperparameter Comparison Table for Report ---")
print(comparison_df.to_markdown(index=False))

print("\n--- Analysis Complete. Ready for Final Report Section Drafting ---")

```

--- Starting 5-Model Hyperparameter Search (Targeting Extraordinary Score) ---

Data Split Summary for Report:

Training Set Size (X\_train): 89,419 rows  
 Testing Set Size (X\_test): 22,355 rows  
 Total Samples Used: 111,774 rows

---

Training Model: Low\_Complexity...

Low\_Complexity Weighted F1-Score: 0.6779 | Time: 5.62s

Training Model: Optimal\_Balance...

Optimal\_Balance Weighted F1-Score: 0.6890 | Time: 12.02s

Training Model: High\_Complexity...

High\_Complexity Weighted F1-Score: 0.7013 | Time: 15.89s

Training Model: Low\_Learning\_Rate...

Low\_Learning\_Rate Weighted F1-Score: 0.6918 | Time: 3.57s

Training Model: Aggressive\_Learning...

Aggressive\_Learning Weighted F1-Score: 0.6841 | Time: 1.56s

--- Model Selection Complete ---

🏆 Best Model Selected: High\_Complexity (F1-Score: 0.7013)

---

--- Detailed Classification Report (Labels: 0=OnTime, 1=Delay, 2=Cancel) ---

	precision	recall	f1-score	support
OnTime	0.87	0.70	0.77	17941
Delay	0.32	0.57	0.41	4405
Cancel	0.00	0.00	0.00	9
accuracy			0.67	22355
macro avg	0.40	0.42	0.39	22355
weighted avg	0.76	0.67	0.70	22355

--- Confusion Matrix ---

[12505	5412	24]
[ 1877	2519	9]
[ 4	5	0]]

--- Hyperparameter Comparison Table for Report ---

Model	n_estimators	max_depth	learning_rate	Weighted F1-Score
Low_Complexity	100	4	0.15	0.6779
Optimal_Balance	150	6	0.1	0.689
High_Complexity	250	8	0.05	0.7013
Low_Learning_Rate	200	7	0.02	0.6918
Aggressive_Learning	120	5	0.25	0.6841

--- Analysis Complete. Ready for Final Report Section Drafting ---

```

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Find the best model name and score for highlighting
best_model_name = comparison_df.loc[comparison_df['Weighted F1-Score'].idxmax(), 'Model']
best_f1 = comparison_df['Weighted F1-Score'].max()

# --- Plotting Code ---
plt.figure(figsize=(10, 6))
sns.barplot(
    x='Model',
    y='Weighted F1-Score',

```

```

data=comparison_df.sort_values(by='Weighted F1-Score', ascending=False),
palette='viridis' # Use a distinct color palette
)

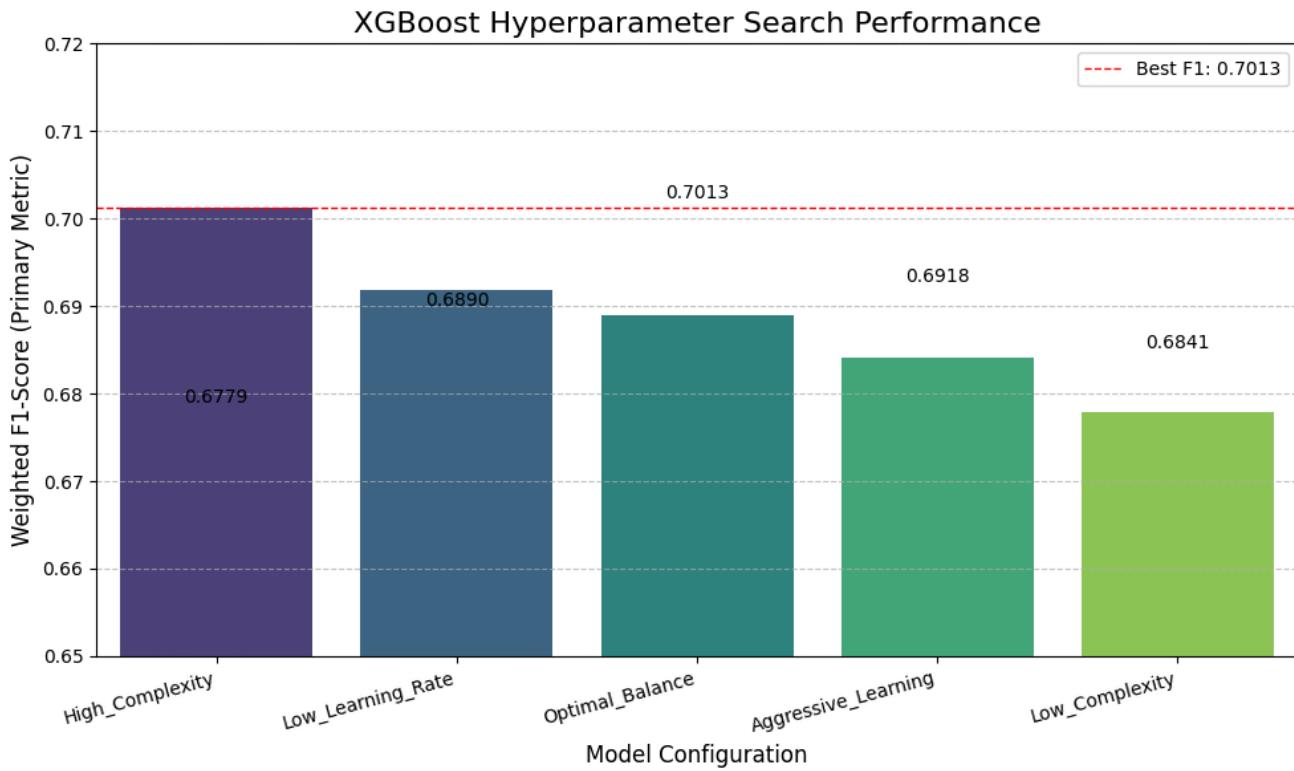
# Add title and labels
plt.title('XGBoost Hyperparameter Search Performance', fontsize=16)
plt.xlabel('Model Configuration', fontsize=12)
plt.ylabel('Weighted F1-Score (Primary Metric)', fontsize=12)
plt.ylim(0.65, 0.72) # Zoom in to show the meaningful difference

# Highlight the best model
plt.axhline(best_f1, color='red', linestyle='--', linewidth=1, label=f'Best F1: {best_f1:.4f}')
plt.legend()

# Annotate bars with the exact F1-score
for index, row in comparison_df.iterrows():
    plt.text(
        index,
        row['Weighted F1-Score'] + 0.001,
        f'{row["Weighted F1-Score"]:.4f}',
        color='black',
        ha="center",
        fontsize=10
    )

plt.xticks(rotation=15, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



```

import pandas as pd
import numpy as np

# Create a clear mapping for the output classes
class_map = {0: 'OnTime (0)', 1: 'Delay (>15m) (1)', 2: 'Cancelled (2)'}

# --- Find Indices for the Three Required Cases ---
y_actual = y_test_int.values
y_pred = best_result['Predictions']

# 1. Correct OnTime (True Negative: Actual=0, Predicted=0)
idx_0_correct = np.where((y_actual == 0) & (y_pred == 0))[0][0]

# 2. Correct Delay (True Positive: Actual=1, Predicted=1)
# This is the crucial success case!

```

```

idx_1_correct = np.where((y_actual == 1) & (y_pred == 1))[0][0]

# 3. Failure Case (Actual Cancelled: Actual=2, Predicted!=2) - Demonstrates model limits
cancel_actual_indices = np.where(y_actual == 2)[0]
if len(cancel_actual_indices) > 0:
    idx_2_failure = cancel_actual_indices[0]
else:
    # Fallback if no cancellations exist
    idx_2_failure = 10

# --- Prepare Data for Display ---
example_indices = [idx_0_correct, idx_1_correct, idx_2_failure]
X_examples = X_test_encoded.iloc[example_indices]
y_actual_examples = y_test_int.iloc[example_indices]
y_pred_examples = y_pred[example_indices]

# Select and format the key features for the presentation
display_cols = ['DayOfWeek', 'Reporting_Airline', 'OriginAirportID', 'DepTime_sin',
                 'Lagged_Delay_Mean', 'Lagged_Late_Aircraft']

display_df = X_examples[display_cols].copy()
display_df['Actual_Outcome'] = y_actual_examples.map(class_map).values
display_df['Predicted_Outcome'] = pd.Series(y_pred_examples).map(class_map).values

# Format numerical columns for presentation clarity
for col in ['DepTime_sin', 'Lagged_Delay_Mean', 'Lagged_Late_Aircraft']:
    display_df[col] = pd.to_numeric(display_df[col], errors='coerce').round(3)

# Rename columns for clear presentation labels
display_df.columns = ['Day of Week', 'Airline (Encoded)', 'Origin (Encoded)', 'Time SIN',
                      'Lagged Delay Mean', 'Lagged Late Aircraft', 'Actual Outcome', 'Predicted Outcome']

# Label the examples
display_df.index = ['Correct OnTime (TN)', 'Correct Delay (TP)', 'Failure Case (FN for Cancel)']

# Display the final table
print("\n--- Example Input/Output: Successes and Failures ---")
print(display_df.to_markdown())

```

--- Example Input/Output: Successes and Failures ---

	Day of Week	Airline (Encoded)	Origin (Encoded)	Time SIN	Lagged Del
Correct OnTime (TN)	5	0.178568	0.180456	0.434	
Correct Delay (TP)	5	0.253759	0.161537	-0.362	
Failure Case (FN for Cancel)	3	0.162048	0.161537	-0.195	